

# UI Hacking: Presentation Creation & Voiceover Patterns

---

Competitive UX research for Designrr — book/manuscript → presentation → voiceover flow

Prepared 2026-06-19 · 20 products analyzed across 4 clusters · Sources: official docs, help centers, G2/Capterra/Trustpilot reviews, and verified Mobbin screen captures

This report reverse-engineers how 20 leading products solve pieces of the same problem Designrr faces: turning long-form or prompt-based content into a structured visual deck, optionally narrated with AI voice. The goal is not to copy any single product, but to extract reusable principles — entry points, generation checkpoints, editing granularity, voice/avatar selection, and progress communication — that can be adapted to Designrr's book-to-presentation-to-voiceover flow.

**Mobbin-verified products** (real indexed screenshots/flows): Gamma, Pitch, Canva, Synthesia, ElevenLabs, Descript, NotebookLM. All others are sourced from official product documentation, help centers, and review aggregators — no dedicated Mobbin coverage was found for them.

## Cluster A — AI-First Presentation Generators

### 1. Gamma Mobbin-verified

**Flow:** Pick Presentation/Webpage/Document/Social → prompt, paste text, or import (PDF/DOCX/URL) → AI drafts an **editable outline** (titles + bullet previews) → user sets card count, text density, style, language → ~30–60s full generation → card-canvas editor with AI chat sidebar.

**Notable decisions:** outline-before-build gate; explicit numeric slide-count control; "shuffle" on example prompts; live character counter; Remix labeled BETA with inline feedback ask.

**Why it works:** the outline is a cheap checkpoint that catches structural mistakes before the expensive visual generation step.

**Drawbacks:** non-linear card layout unfamiliar to PowerPoint users; no native audio export.

**Patterns:** outline checkpoint · parallel entry ramps (prompt/paste/import) · explicit pre-generation controls.

Refs: [gamma.app/products/presentations](https://gamma.app/products/presentations) · Mobbin: [outline review screen](#), [paste-text screen](#)

### 2. Tome discontinued / cautionary

**Flow (historical):** single prompt → full deck in <60s, no outline step, non-linear "tile" narrative.

**Why it briefly worked:** skipping the outline maximized perceived speed/magic for demos.

**Drawback — important data point:** the no-outline, one-shot approach gave users less structural control, and the standalone product was discontinued (pivoted away from presentations, March 2025) — generation speed alone wasn't a durable differentiator.

**Pattern (cautionary):** one-shot generation without a review step trades control for speed — risky as a sole differentiator.

Refs: [Tome AI 2025 pivot deep-dive](#)

### 3. Beautiful.ai

**Flow:** "DesignerBot" prompt or Smart Slide template → AI drafts structure/copy/imagery → **layout guardrails enforced live** (add a bullet → auto-reflow; swap image → auto-resize).

**Why it works:** removes the *possibility* of an ugly slide rather than just nudging toward good defaults — directly answers "I'm not a designer."

**Drawbacks:** "no override" frustrates power users; weak at AI-generated diagrams.

**Patterns:** constraints baked into the object model, not just suggested · branded AI persona as the face of generation.

Refs: [beautiful.ai/ai-presentations](https://beautiful.ai/ai-presentations)

#### 4. Pitch Mobbin-verified

**Flow:** Home offers "Start new," "Start with AI," "Import" as equal-weight tiles → AI prompt generates a deck with a side panel where **Prompt/Fonts/Colors are independently re-rollable** → post-generation welcome modal: "explore myself" vs "quick tour" → slide-level analytics (visitor feed, view-over-time) live inside the editor.

**Why it works:** granular re-roll avoids all-or-nothing regeneration frustration; opt-in marketing toggle defaults to off, building trust.

**Drawbacks:** generation quality more basic than Gamma's.

**Patterns:** facet-level regeneration · privacy-respecting default toggles · post-output tour-or-explore choice.

Refs: Mobbin [Start-with-AI flow](#), [editor tour screen](#)

#### 5. Plus AI (Slides/PowerPoint add-in)

**Flow:** Lives inside Slides/PPT as a native panel → prompt → editable outline → ~10-slide deck as native objects in <2 min → per-slide "ask AI to rewrite" scoped to one selected slide. Separate "Presentation Narrator" add-on bolts on TTS/avatar narration.

**Why it works:** staying inside the tool people already trust removes the #1 objection ("will this match my template?").

**Drawbacks:** weaker novel visual layouts than dedicated design tools.

**Patterns:** embed AI inside the existing trusted workspace · scope AI edits to one selected element · keep narration a separate, optional layer.

Refs: [plusai.com](https://plusai.com)

## Cluster B — Template-Based Design Tools

#### 6. Canva Mobbin-verified

**Flow:** Home shows Your designs/Templates/Canva AI as equal tabs → search-first template gallery with style chips, or AI chat (with quick actions + doc attach), or PPTX import → Present mode splits into Standard/Autoplay/Presenter/"**Present and record.**"

**Why it works:** search-first beats folder-browsing for decision speed; treating "record narration" as a mode of Present (not a buried feature) makes voiceover discoverable.

**Drawbacks:** PPTX export frequently breaks layout/fonts; free tier capped at 10 AI generations/month.

**Patterns:** AI-prompt and template-browse as equal-weight tabs · narration as a Present-mode option, not a separate menu.

Refs: [canva.com/magic-design](https://canva.com/magic-design) · Mobbin: [Create-a-design flow](#), [Canva AI flow](#)

## 7. Visme

**Flow:** Chat-driven AI with doc/URL upload buttons **inside the prompt composer** → style suggestions → draft → editor with per-element AI Text-to-Speech (6 voices, tone/speed, 4000-char limit).

**Why it works:** voiceover scoped to specific blocks/hotspots matches async self-review use cases (e.g., a prospect reviewing a sales deck alone).

**Drawbacks:** free plan lets you build a full deck then blocks download — a bait-and-switch reviewers flag.

**Patterns:** "upload doc/URL" lives inside the AI prompt box itself · voiceover scoped to individual content blocks, not just whole-deck.

## 8. Prezi

**Flow:** Blank/template/AI (prompt or PDF/PPT/Word upload) → zoomable, non-linear "path" presentation → narration via right-click "Add Sound to Path Step" (per-object) or full Prezi Video record with a **"content only" no-webcam toggle**.

**Why it works:** the no-camera toggle directly removes recording anxiety, a real dropout point.

**Drawbacks:** zoom/pan navigation can disorient viewers in long decks; export to PPT flattens the zoom.

**Patterns:** no-camera narration option as a first-class toggle · audio attachable to individual canvas objects, not just deck-level.

## 9. Slidesgo / SlidesAI

**Flow:** Slidesgo = searchable template marketplace (free tier = mandatory attribution slide). SlidesAI = Google Slides sidebar: prompt **or** "Create from my Text" (paste up to 2,500 chars) → AI shows an **outline for approval** → generates slides with images, in-place.

**Why it works:** outline-approval mirrors validating a table of contents before writing — cheap correction point.

**Drawbacks:** 2,500-char cap forces manuscript content to be trimmed before conversion; billing/cancellation complaints.

**Patterns:** distinct "describe a topic" vs "paste my existing text" modes · outline approval before full generation.

## 10. Renderforest

**Flow:** Template gallery → **"Select Scenes"** step (choose/reorder from hundreds of scenes before adding content) → drag-drop media → AI voiceover applied per scene (language/voice chosen at assembly time) → rendered as an actual video file with a resolution choice.

**Why it works:** structure-before-content reduces blank-canvas overwhelm; framing output as "a video to render" sets accurate expectations.

**Drawbacks:** confusing two-currency pricing (subscription + credits); heavy free-tier watermark.

**Patterns:** separate "arrange structure" from "fill content" as two sequential steps · apply voice/language at the scene-assembly stage, not as a final pass.

## Cluster C — AI Avatar Video Generators

### 11. Synthesia Mobbin-verified

**Flow:** Blank/template/import (PPTX/PDF/Word/Slides, up to 150 slides) → each slide becomes a **scene**: speaker notes auto-populate the script box → pick avatar from a searchable grid (text-prompt search) → voice/language picker with inline play-to-preview → "Generate" (top-right, distinct from autosave).

**Why it works:** scene = slide mental model leverages existing PowerPoint literacy; inline preview kills commitment anxiety before burning render credits.

**Drawbacks:** 8–12 min re-render per change is a real iteration tax; lip-sync uncanny valley.

**Patterns:** scene as atomic unit (script+avatar+visual) · script editor paired with inline timeline · preview-before-commit on every voice/avatar pick · visually separate draft-state from render-commit action.

Refs: [docs.synthesia.io/docs/powerpoint](https://docs.synthesia.io/docs/powerpoint) · Mobbin: [scene script editor w/ timeline](#), [voice/language picker](#), [generate modal](#)

### 12. HeyGen

**Flow:** Script editor central → avatar + "Avatar Engine" mode (Unlimited/IV) with **expressiveness slider + motion-prompt field** kept separate from the script text → voice library with faceted filters (region/accent/age/emotion) + favorites → bulk "apply this voice to all scenes."

**Why it works:** faceted filtering matches how people search large catalogs (700+ avatars); splitting "what to say" from "how to perform it" gives control without exposing animation complexity.

**Drawbacks:** every re-render burns credits even on failures — the top complaint (Trustpilot 2.4/5 despite G2 4.8/5, signaling billing-driven dissatisfaction).

**Patterns:** faceted filtering for large libraries · saved favorites · decouple script from performance/expressiveness controls · bulk-apply-to-all-scenes shortcut.

### 13. Elai.io

**Flow:** Blank/AI topic-to-script/import PPTX/PDF → storyboard of scenes (up to 100) → avatar types as tabs (selfie/studio/photo/mascot) + voice cloning as a peer option → AI script drafting embedded directly in the script box → renders **without a full preview**.

**Why it works:** typed avatar categories cut decision fatigue vs. scrolling 80 lookalikes.

**Drawbacks — anti-pattern:** no full preview before render consumes a limited monthly minute quota; 2–3 revision rounds can exhaust it.

**Patterns:** embed AI drafting in the script field itself · categorize avatars by tone/use-case, not just demographics. **Anti-pattern:** never skip a pre-render preview when render time is metered.

## 14. Colossyan

**Flow:** "Cora" agent takes doc/URL/script → **staged approval: approach** → **plan** → **render**, each requiring explicit user confirmation before the next (expensive) step → up to 4 avatars per scene for dialogue/training content → translation re-syncs text+audio+lip-sync as one atomic action across 100+ languages.

**Why it works:** staged approval directly targets the "wasted render minutes" pain reported across the whole category.

**Drawbacks:** still no preview before final export; 10–15+ min processing per clip.

**Patterns:** gate expensive compute behind staged approval (intent→plan→render) · multi-speaker per scene · treat translation as one atomic multi-asset operation.

## 15. Powtoon

**Flow:** Blank/template/"Imagine Script AI"/"Anything-to-Video" (upload doc, pick output style: animated/presentation/footage-based) → Stage+Timeline editor with **object-level** (not scene-level) duration bars → dual preview modes (Video vs Slideshow) before export.

**Why it works:** object-level timeline appeals to users wanting fine animation control; choosing output style upfront avoids late surprises.

**Drawbacks:** HD/no-watermark gated to pricier plans; reported export glitches.

**Patterns:** expose object-level timing on animation content · choose output style/tone before generation · multiple preview modes as a final gate before costly export.

## Cluster D — Voiceover, TTS & Document-to-Audio Tools

### 16. Murf.ai

**Flow:** Paste/type script → auto-split into sentence blocks → 200+ voice library filtered by language/accent/gender/age/style → per-word emphasis markup, timed pauses, **persistent pronunciation dictionary** → generate → **regenerate a single block** without re-rendering the rest → timeline for music/SFX layering.

**Why it works:** addressable blocks turn "generate everything" into low-risk, granular edits.

**Drawbacks:** non-English voices noticeably less natural; strict 24hr/10-min refund policy is the top complaint.

**Patterns:** faceted voice filters · block-level regeneration · persistent pronunciation overrides that carry across projects.

### 17. Descript (Overdub/Regenerate) Mobbin-verified

**Flow:** Auto-transcribe audio/video into an editable script → **editing the text re-cuts the media** → select a passage → contextual floating toolbar ("Regenerate," "Edit for clarity," "Remove filler words") → regenerating unchanged words re-synthesizes tone/delivery; regenerating **new** words requires Overdub voice-cloning consent first.

**Why it works:** treating audio editing as text editing collapses a technical skill into a familiar one; contextual actions keep the UI clean until intent is signaled.

**Drawbacks:** quality degrades on long passages; AI credits deplete fast.

**Patterns:** text-edit = media-edit · contextual selection toolbar · separate consent tiers for "same words, new delivery" vs "new words, new speech."

Refs: Mobbin [Editing Script Clarity flow](#)

### 18. ElevenLabs (Studio + Reader) Mobbin-verified

**Flow:** Import doc → auto-split into **Chapters** (sidebar list) → each chapter is independently converted with a status badge ("Converted") and its own **version history** → voice-reuse panel shows "Used in Project" for one-click consistency across chapters. Reader app: paste text/PDF/ePub/URL → instant narration with word-by-word highlighting.

**Why it works:** chapter-level status turns a long job into a visible checklist, killing "is this done?" anxiety; versioning (not overwriting) protects against destructive AI regenerations.

**Drawbacks:** credit-based pricing — unused credits vanish on cancellation.

**Patterns:** decompose long documents into independently-convertible chapters with status indicators · keep version history per section instead of overwriting · surface "voice already used here" for reuse.

Refs: Mobbin [Creating a Project flow](#), [Voice Library screen](#)

## 19. Google NotebookLM — Audio Overview Mobbin-verified

**Flow:** Upload sources → Studio panel → "Audio Overview" → **pre-generation modal:** format presets (Deep Dive/Brief/Critique/Debate) + free-text "Customize" steering box → Generate (async, user can keep working) → audio is one card among other Studio outputs (study guide, FAQ, mind map) → "Interactive mode" lets the user verbally barge into the synthetic conversation.

**Why it works:** two-host podcast format exploits parasocial familiarity, making dense source material feel low-effort; front-loading steering avoids mid-generation decision paralysis.

**Drawbacks — the single most-cited limitation: no post-generation editing** — any mispronunciation or wrong emphasis requires regenerating the *entire* file. This is the clearest cautionary example for Designrr.

**Patterns:** when output can't be edited after, concentrate all steering into one pre-generation step · named output-style presets instead of raw parameters · async non-blocking generation. **Anti-pattern: no segment-level fix-up.**

Refs: Mobbin [Generating an Audio Overview flow](#)

## 20. Lumen5 (blog/doc-to-video)

**Flow:** Paste blog URL/text → AI splits into sentence-level scenes, matches stock footage, drafts voiceover → Script Composer (text editing separate from visual editing) → card-based storyboard, drag-and-drop reorder/retime → optional custom voiceover upload instead of built-in TTS.

**Why it works:** 1:1 sentence-to-scene mapping gives an immediately legible structure — text becomes video structure.

**Drawbacks:** generic/repetitive stock footage is the top complaint; built-in TTS still sounds robotic.

**Patterns:** map source-text structure 1:1 onto output units · keep script-editing and visual-editing surfaces separate · always allow manual override of any AI-automated step.

# 1. Pattern Inventory

| Pattern                                                 | Products                               | Why it works                                                         | Risks                                            | Recommendation                                                                                                         |
|---------------------------------------------------------|----------------------------------------|----------------------------------------------------------------------|--------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|
| Outline-before-build checkpoint                         | Gamma, Plus AI, SlidesAI, Beautiful.ai | Cheap correction point before expensive generation                   | Extra step can feel slow to impatient users      | Show a chapter/section outline pulled from the manuscript; let users pick/reorder sections before slides are generated |
| Parallel entry ramps (prompt/paste/import/blank)        | Gamma, Canva, Visme, Prezi             | Meets users wherever their content already lives                     | Too many choices → paralysis on the first screen | Cap at 3 tiles: "From a book," "From a prompt/topic," "Start blank"                                                    |
| Granular, scoped regeneration (facet/block/scene-level) | Pitch, Murf, Descript, ElevenLabs      | Avoids all-or-nothing redo; preserves good parts                     | More UI surface to design/discover               | Critical for Designrr: regenerate one slide's narration without re-rendering the book                                  |
| Scene = atomic unit (text + visual + voice)             | Synthesia, HeyGen, Elai, Colossyan     | Matches existing slide-deck mental model; supports incremental edits | Can fragment tone consistency across scenes      | Model each presentation page as a scene object carrying text, voice, and timing together                               |
| Preview-before-commit for voice/avatar                  | Synthesia, HeyGen, Murf                | Removes regret/risk before spending render credits                   | None significant                                 | Always show an inline play icon in the voice picker before applying                                                    |
| Staged approval before expensive compute                | Colossyan, Gamma's outline gate        | Prevents wasted compute, builds trust                                | More clicks before output                        | Confirm scope ("narrate all 12 chapters?") before committing to a full-book render                                     |
| Faceted filtering for large libraries                   | HeyGen, Murf, ElevenLabs, Canva        | Scales to hundreds of options; matches search behavior               | Too many facets overwhelms casual users          | Filter voices by language/accent/tone; default to a short "recommended" list                                           |
| Persistent personalization (pronunciation, favorites)   | Murf, ElevenLabs, HeyGen               | Reduces repetitive correction over a long project                    | Needs account-level state                        | Save a preferred narrator + custom pronunciations per book project                                                     |
| Chapter decomposition with status tracking              | ElevenLabs, NotebookLM, Synthesia      | Turns a long job into a visible checklist                            | More visual chrome; needs an aggregate view too  | Show per-chapter "Pending/Converted/Regenerate" badges for book-length voiceover                                       |
| Version history instead of destructive overwrite        | ElevenLabs                             | Protects against AI regressions; builds trust                        | Storage/complexity cost                          | Keep the last few narration takes per slide; let users revert                                                          |
| Embed AI inside the existing/trusted workspace          | Plus AI, SlidesAI                      | Removes compatibility/migration objection                            | N/A — Designrr already owns the workspace        | Keep book→presentation conversion in-place, not an export-then-reimport detour                                         |
| Design guardrails baked into the object model           | Beautiful.ai                           | Removes the possibility of a bad layout, not just a nudge            | Frustrates power users with "no override"        | Auto-paginate/resize when pasted book text overflows a slide                                                           |
| Text-edit = media-edit                                  | Descript                               | Reuses a familiar skill (word processing) for a                      | Implementation complexity (forced                | Editing slide text should flag stale narration as "out of sync —                                                       |

|                                                                |                       |                                                            |                                                               |                                                                                  |
|----------------------------------------------------------------|-----------------------|------------------------------------------------------------|---------------------------------------------------------------|----------------------------------------------------------------------------------|
|                                                                |                       | technical one (audio editing)                              | alignment                                                     | regenerate?"                                                                     |
| No-camera / audio-only narration toggle                        | Prezi                 | Removes recording anxiety                                  | Minor                                                         | Default to audio-only; never force an avatar/webcam                              |
| Output tone/style chosen upfront                               | Powtoon, Renderforest | Sets expectations before generation; avoids late surprises | Minor                                                         | Ask narration tone (professional/conversational/dramatic) before generating      |
| Async, non-blocking long-running generation                    | NotebookLM, Synthesia | User isn't trapped waiting; can keep working               | Needs a clear completion notification                         | Let users keep editing while a full-book voiceover renders; notify on completion |
| Pre-generation customization when output can't be edited after | NotebookLM            | Front-loads control to match a real constraint             | Worse than granular post-editing wherever editing IS possible | Avoid this constraint; favor Descript/Murf-style post-edit instead               |
| Default privacy-friendly toggle state (opt-in off)             | Pitch                 | Builds trust, avoids dark-pattern perception               | None                                                          | Apply to any marketing/data-sharing toggles                                      |
| 1:1 mapping of source structure to output structure            | Lumen5, ElevenLabs    | Legible mental model — no guessing where content "lives"   | Rigid mapping may not fit unusual book structures             | Map chapters/headings to slides by default; allow manual re-split                |

## 2. Common Trends

- **Outline/plan-before-generate is nearly universal** among AI-first tools (Gamma, Plus AI, SlidesAI, Colossyan) — the standard mitigation for one-shot generation's unpredictability.
- **Granular regeneration is the single biggest differentiator** between loved and hated AI tools. Every product lacking it (NotebookLM Audio Overview, Elai, one-shot-era Tome) draws complaints about wasted time/credits on full re-runs for one bad line.
- **Voice/avatar selection always pairs with inline preview** — no product makes you commit blind.
- **Long jobs get decomposed into visible units** (scenes, chapters, Studio cards) with status — never a single opaque spinner for multi-minute generation.
- **Render/credit economics shape the UX directly**: staged approvals, previews, and granular edits exist specifically because generation costs real money/compute — exactly the situation Designrr will be in with book-length voiceover.
- **Multiple input modalities are offered in parallel**, not gated behind a single forced path — prompt, paste, import, and blank coexist as equal-weight choices.

## 3. Best Practices to Adopt

1. Outline-first checkpoint before slide/voice generation.
2. Granular, scoped regeneration (per slide/paragraph, never whole-book).
3. Scene-as-atomic-unit bundling text + voice + visual.
4. Preview-before-commit on every voice/avatar choice.
5. Chapter-level decomposition with visible status for long-running jobs.
6. Text-edit triggers media-edit (flag stale narration on text change).

7. Persistent personalization (preferred voice, pronunciation dictionary) per project.
8. Multiple parallel entry points matching where the user's content already lives (book vs. prompt vs. blank).

## 4. Anti-Patterns to Avoid

---

1. **One-shot generation with no outline and no post-edit** (Tome's failed model; NotebookLM Audio Overview's top complaint) — pick one safety net or the other, ideally both.
2. **No preview before consuming metered render credits** (Elai, Colossyan).
3. **All-or-nothing regeneration** forcing a full re-render to fix one line.
4. **Opaque/punitive credit economics** (HeyGen's Trustpilot 2.4 despite G2 4.8 — billing, not quality, drives churn).
5. **Build-then-paywall bait-and-switch** — letting users complete a full project, then blocking export/download (Visme, Slidesgo free tiers).
6. **Rigid guardrails with zero override** (Beautiful.ai's "no override" complaint) — fine for guardrails, not fine for blocking power users entirely.
7. **Spatial/non-linear navigation for inherently linear content** (Prezi's disorientation/motion-sickness complaints) — a book is sequential; don't fight that.

## 5. Design Recommendations for Designrr

---

1. **Three entry tiles, not a sequential wizard:** "From a book/manuscript" (auto-pulls chapter outline), "From a prompt/topic," "Start blank." Equal weight, same screen — don't force book users down a different path than prompt-first users.
2. **Auto-generate an editable chapter/section outline** when starting from a manuscript. Let the user select which chapters become slides, merge/split sections, reorder — before any slide design or voice work happens.
3. **Model each slide as a "scene" object:** layout + text + voiceover script + assigned voice, bundled but independently editable. This directly supports "voiceover can be added at some point, not necessarily upfront" — voiceover becomes a property attached to a scene whenever, not a separate top-level pass.
4. **Never force whole-book commitment for voiceover.** Offer per-slide or per-chapter generation with status badges (Pending/Converted/Regenerate), mirroring ElevenLabs' Chapters pattern — highest-leverage adoption given book-length content.
5. **Granular regeneration is non-negotiable.** A user editing chapter 7 of a 20-chapter book should regenerate only that narration, never the whole book (Murf/Descript pattern).
6. **Always preview voice samples inline** before applying — to a slide or to the whole book.
7. **Make full-book voiceover generation async and non-blocking**, with a completion notification — books are long; don't trap the user on a spinner.
8. **Persist a preferred narrator voice + pronunciation list per book project** so recurring character/brand names don't need re-correction every chapter.
9. **Gate full-book operations behind a lightweight scope-confirmation step** ("Generate voiceover for all 18 chapters — estimated X minutes/credits?") before committing compute — borrowed from Colossyan's staged approval.
10. **Default any data-sharing/marketing toggles to off**, and never block final export after a user has invested effort building (avoid the Slidesgo/Visme bait-and-switch) — protects trust in an existing paid funnel.
11. **Give the no-book/AI-first path equal polish:** prompt → outline-checkpoint → deck, same quality bar as the book-derived flow, so "presentation first, book later" users aren't second-class.
12. **Stick to a linear scene-rail/slide-sequence model**, not a spatial/zooming metaphor — matches both a book's inherent linearity and the voiceover-per-scene pattern used by every avatar-video competitor.

13. **Flag stale narration on text edits** ("Slide text changed — regenerate narration?") rather than letting script and audio silently drift apart, borrowing Descript's text-edit-is-media-edit philosophy.

## 6. Screenshots & References

---

Full source list is embedded inline above per product. Real, verified screenshots/flows exist on Mobbin for: **Gamma, Pitch, Canva, Synthesia, ElevenLabs, Descript, NotebookLM**. For **Tome, Beautiful.ai, Plus AI, Visme, Prezi, Slidesgo/SlidesAI, Renderforest, HeyGen, Elai.io, Colossyan, Powtoon, Murf.ai, Lumen5** — no dedicated Mobbin coverage was found; those are sourced from official product docs, help centers, and G2/Capterra/Trustpilot reviews (linked inline above).